

Abstract

The exponential expansion of the e-commerce sector has precipitated a logistical crisis, specifically within the "last mile" of delivery, the final segment of the supply chain responsible for transporting goods from distribution hubs to end consumers. This segment is surely inefficient and costly, often responsible for more than half of total shipping expenses. The core computational challenge underlying this inefficiency is the Vehicle Routing Problem (VRP), an NP hard combinatorial optimization problem. While traditional metaheuristics such as Ant Colony Optimization (ACO) have provided functional solutions for decades, they increasingly struggle to handle the scale and complexity of modern urban delivery networks, often converging prematurely to suboptimal routes.

This research proposes a novel Hybrid ACO framework that integrates Machine Learning (ML) models specifically Artificial Neural Networks (ANN), K-Nearest Neighbors (KNN), and Long Short Term Memory (LSTM) networks as dynamic heuristic guides for the optimization swarm. Leveraging the LaDe (Last mile Delivery) dataset, a massive repository of over 1.4 million real world delivery trajectories provided by Cainiao AI, this study conducts a comprehensive comparative analysis of these hybrid architectures.

The experimental results, derived from a GPU accelerated simulation environment, reveal a significant divergence in performance among the tested models. The KNN-ACO approach emerged as the superior architecture, achieving a best route distance of 2.46 km on the test set. This represents a 73.47% improvement over stochastic initialization and outperforms established baselines such as TimeGreedy (2.15 km) and approximates the performance of state of the art Graph Neural Networks. In contrast, the LSTM-ACO model, despite its theoretical capacity for sequence learning, failed to capture the spatial topology of the delivery network, resulting in a significantly poorer performance of 7.88 km. These findings suggest that for static routing problems characterized by spatial clustering, memory based learning mechanisms like KNN provide a more robust and computationally efficient guidance system than complex deep learning architectures.

1. Introduction

1.1 Background and Context

In the modern digital economy, the efficiency of logistics networks is a primary determinant of commercial success. The phenomenon of "last mile delivery" has evolved from a purely operational concern into a strategic differentiator. With the advent of same day delivery promises and the explosion of online retail volumes, logistics providers face the daunting task of routing thousands of vehicles to millions of disparate locations daily. This logistical challenge is compounded by urbanization, traffic congestion, and the customer's expectation of precise delivery windows.

The mathematical formulation of this operational challenge is the Vehicle Routing Problem (VRP). The VRP asks, "What is the optimal set of routes for a fleet of vehicles to traverse in order to deliver to a given set of customers?" As the number of delivery points increases, the number of possible route permutations grows factorially, rendering exact mathematical solutions impossible for large scale instances. Consequently, the industry has historically relied on heuristics rules of thumb and metaheuristics, which are higher level procedures designed to find sufficiently good solutions in a reasonable timeframe.

Among metaheuristics, Ant Colony Optimization (ACO) has held a prominent position. Inspired by the foraging behaviour of biological ants, ACO utilizes a population of agents that construct solutions probabilistically. Agents leave a "pheromone" trail on successful paths, increasing the probability that subsequent agents will follow the same route. While ACO is robust, it is inherently "blind" to the complex structural patterns of a city's geography. It relies on simple distance calculations and iterative trial and error, often requiring thousands of iterations to find a solution that a human planner might spot intuitively.

1.2 Motivation and Research Gap

The motivation for this study stems from the rapid advancements in Artificial Intelligence (AI) and Machine Learning (ML). While Deep Learning has revolutionized fields such as computer vision and natural language processing, its application to combinatorial optimization, termed Neural Combinatorial Optimization (NCO) is still maturing.

Current literature reveals a dichotomy, purely heuristic methods are fast but often suboptimal, while end to end deep learning methods require massive training datasets and can be brittle when problem sizes change. There is a significant gap in the research regarding "Hybrid Intelligence" systems that combine the robust search capabilities of metaheuristics with the pattern recognition capabilities of machine learning. specifically, there is limited understanding of which type of learning spatial (KNN), non linear approximation (ANN), or sequential (LSTM) is best suited to augment swarm intelligence in a logistics context.

1.3 Research Question

The primary research question addressed by this study is, to what extent can Machine Learning augmented heuristic guidance improve the convergence speed and solution quality of Ant Colony Optimization in large scale last mile delivery scenarios?

Sub questions include:

1. Does the incorporation of local spatial memory (via KNN) outperform global function approximation (via ANN) in guiding optimization agents?
2. Is the Vehicle Routing Problem sufficiently "sequential" for Recurrent Neural Networks (like LSTM) to provide meaningful predictive value compared to spatial approaches?
3. Can hybrid algorithms be efficiently scaled to real world datasets using GPU acceleration?

1.4 Research Objectives

The objectives of this research are fourfold. First, the study aims to analyse and characterise the spatial topology of the LaDe dataset in order to identify clustering and connectivity patterns relevant to last mile delivery. Second, a modular hybrid optimization framework is developed to enable the integration of multiple machine learning guidance models with Ant Colony Optimization. Third, the proposed hybrid methods are systematically evaluated using route length and convergence behaviour as primary performance indicators. Finally, the study seeks to provide theoretical insight into the suitability of different learning paradigms for spatial routing problems.

2. Related Work

2.1 The Evolution of VRP Solvers

The Vehicle Routing Problem (VRP) has been a central focus of Operations Research since its formulation by Dantzig and Ramser (1959). Early approaches utilized exact methods such as Branch and Cut, which guarantee optimality but are computationally prohibitive for problems with more than a few hundred nodes. This limitation necessitated the development of heuristics, starting with the Clarke Wright Savings algorithm (1964), which constructs routes by iteratively merging sub routes to maximize cost savings.

In the 1990s, the focus shifted to metaheuristics. Genetic Algorithms (GA), Tabu Search, and Simulated Annealing became the standard. Ant Colony Optimization (ACO), introduced by Dorigo (1992), offered a unique advantage for routing problems, its constructive nature allows it to handle constraints (like capacity or time windows) more naturally than perturbative methods like Genetic Algorithms. Bell and McMullen (2004) empirically demonstrated that ACO outperforms GA in constrained VRP scenarios, attributing this to the pheromone mechanism's ability to retain a "memory" of good sub structures within the graph.

2.2 Neural Combinatorial Optimization

The application of neural networks to optimization problems, known as Neural Combinatorial Optimization (NCO), is a burgeoning field. Vinyals et al. (2015) introduced Pointer Networks, a variation of Recurrent Neural Networks (RNNs) capable of outputting a permutation of input indices, effectively "pointing" to the next city in a tour. This was a breakthrough because it allowed a neural network to output a solution with a variable length, a requirement for routing problems.

Building on this, Kool et al. (2019) proposed an Attention Model approach, replacing the RNN with a Transformer architecture. Their model, which forms the basis of the Graph2Route baseline referenced in this study, utilizes self attention mechanisms to understand the relationships between all nodes in the graph simultaneously, rather than sequentially. While these end to end learning models are powerful, they often suffer from generalization issues; a model trained on 50 node problems may fail spectacularly when tested on 100 node problems.

2.3 Hybrid Metaheuristics

Recognizing the limitations of pure learning approaches, researchers have begun exploring hybrids. A survey by Karagul et al. (2024) categorizes these hybrids into two types, those where ML tunes the parameters of the metaheuristic (e.g., adjusting the pheromone evaporation rate using Reinforcement Learning), and those where ML replaces a component of the metaheuristic.

The latter category is the focus of this research. Ahmed et al. (2023) explored using K-Nearest Neighbour's (KNN) to prune the search space for ACO. By only allowing ants to visit the k nearest neighbours of their current location, they significantly reduced computational time. However, this "hard pruning" risks excluding optimal edges that happen to be long (e.g., connecting two distant clusters). This project builds upon this by using ML not to prune the space, but to *bias* the probabilistic choice, allowing for a softer, more flexible integration of learned knowledge.

2.4 The LaDe Dataset Context

Most academic research in VRP relies on synthetic datasets (such as the Solomon instances). However, real world logistics data is significantly messier, exhibiting high spatial clustering and temporal patterns. The LaDe dataset, introduced by Wang et al. (2023), addresses this by providing real trajectories from Cainiao's network in China. The authors of the LaDe dataset demonstrated that simple greedy heuristics often fail on this data because they do not account for the driver's implicit knowledge of the road network knowledge that this project aims to capture using Machine Learning.

3. Research Methodology

3.1 Dataset Selection and Processing

This research utilizes the LaDe (Last mile Delivery) dataset, specifically the subset covering the city of Shanghai (delivery_sh.csv). This subset was chosen for its high density of orders and complex urban topology.

- **Data Volume:** The raw dataset comprises over 1.4 million records.
- **Preprocessing:** The data underwent rigorous cleaning. Records with missing GPS coordinates were removed. A stratified sampling technique was used to extract a subset of 300,000 records for connectivity analysis to ensure computational feasibility while maintaining statistical significance.
- **Coordinate Normalization:** To facilitate the training of Neural Networks, the latitude and longitude values were normalized. This is a crucial step in ML pipelines to prevent numerical instability during gradient descent optimization.

3.2 Spatial Analysis Methodology

Before optimization, it is critical to understand the environment. Connectivity analysis spatial graph was constructed where delivery locations act as nodes. Edges were established between nodes if they fell within a specific geographic radius (0.01 degrees). This mimics the operational reality where a driver services a local "patch" of deliveries. Cluster Identification, the DBSCAN (Density Based Spatial Clustering of Applications with Noise) algorithm was employed. Unlike K-Means, DBSCAN does not require specifying the number of clusters a priori and can identify arbitrarily shaped clusters, making it ideal for detecting neighbourhood's and delivery zones in a city layout.

3.3 Hybrid Algorithm Methodology

The core methodology involves modifying the decision making process of the artificial ants.

3.3.1 Standard ACO Logic

In standard ACO, an ant at location "i" chooses to move to location "j" based on two factors:

1. **Pheromone:** How many ants have successfully used this path before?
2. **Heuristic:** How close is location "j" ? (Inverse of distance).

3.3.2 The Hybrid Extension

This research introduces a third factor, Machine Learning Guidance, the probability of moving to location "j" is now also weighted by a prediction from an ML model. ANN-ACO uses a Feed Forward Neural Network to predict the "quality" of a move. The network takes the coordinates of the current and potential next location and outputs a score. KNN-ACO uses a "memory bank" of successful route segments. If the move from "i" to "j" is spatially similar to a move that was part of a high quality route in the past, the KNN increases the

probability of selection. LSTM-ACO uses a Recurrent Neural Network. The LSTM analyses the sequence of the last 5 locations visited by the ant. It attempts to predict the most logical next step based on the trajectory's history, modelling the "momentum" of the route.

3.3.3 The Co-Evolutionary Loop

The system uses an iterative feedback loop.

1. **Step 1:** The ML model provides guidance to the ants.
2. **Step 2:** The ants construct routes based on this guidance (plus pheromones and distance).
3. **Step 3:** The best route found by the swarm is identified.
4. **Step 4:** This "best route" is used as a training example to update the ML model.

4. Design Specification

4.1 System Architecture

The software is architected as a modular pipeline, ensuring separation of concerns between data handling, optimization logic, and visualization. The proposed system architecture consists of four main components. The data layer is responsible for ingestion, cleaning, and transformation of the raw delivery trajectories. The optimization engine implements the Ant Colony Optimization logic and manages pheromone dynamics and ant populations. Predictor interfaces provide a modular mechanism for integrating ANN, KNN, and LSTM guidance models without altering the core optimization code. Finally, the visualization module generates convergence plots and route visualisations to facilitate qualitative analysis of the optimization process.

4.2 Machine Learning Model Design

Artificial Neural Network (ANN) Designed as a multi layer perceptron. It features an input layer accepting coordinate pairs, two hidden layers with Rectified Linear Unit (ReLU) activations to capture non linear relationships, and a Sigmoid output layer to produce a probability score between 0 and 1. **K-Nearest neighbours (KNN)** Configured to store the vectors of successful route segments. It uses a KD Tree or Ball Tree data structure for efficient querying of nearest neighbours in high dimensional space. The parameter "K" is set to 5 to balance local specificity with noise reduction. **Long Short Term Memory (LSTM)** Designed to handle sequential data. The input is a tensor representing the last 5 coordinates. The LSTM layer maintains an internal cell state to remember long term dependencies, followed by a dense output layer to predict the score of the next candidate location.

4.3 Computational Strategy

Given the computational intensity of calculating probabilities for multiple ants across a fully connected graph, CPU execution is insufficient. The design mandates the use of GPU Acceleration. The system utilizes the CuPy library, which provides a NumPy compatible interface for CUDA. This allows the matrix operations, such as calculating the Euclidean distance between all pairs of nodes to be parallelized on the GPU, resulting in orders of magnitude speed improvements.

5. Implementation

5.1 Data Ingestion and Cleaning

The implementation begins with loading the `delivery_sh.csv` file. A critical implementation detail was the handling of "drift" data GPS points that drift to (0,0) due to sensor error. A filter was implemented to remove any coordinates falling outside the bounding box of the city of Shanghai. The `preprocess_trajectory_data` function handles the normalization. Coordinate values, originally in degrees (e.g., 121.5), were scaled. This is vital for the Neural Networks, as inputs with large variance can cause the gradients to explode or vanish during backpropagation.

5.2 Implementing the Hybrid Optimizers

The Hybrid ANN-ACO, Hybrid KNN-ACO, and Hybrid LSTM-ACO classes were implemented to encapsulate the specific logic of each approach. The KNN Memory Bank in the train KNN method, the system extracts every segment (edge) from the best route found in the current epoch. These segments are pushed to a training buffer. To prevent memory leaks during long runs, a sliding window approach was implemented, keeping only the most recent 1,000 successful segments. The LSTM Sequence Generator, implementing the LSTM required careful handling of data structures. As an ant builds a tour, its history grows. The implementation maintains a rolling buffer of the last 5 visited nodes for every single ant in the colony. This buffer is fed into the LSTM at every step to generate the guidance value for the next decision.

5.3 Optimization Loop

The `optimize` method runs for a user defined number of epochs (set to 100 for the experiments). Inside this loop, Guidance Generation, the ML model predicts a "desirability matrix" for the entire graph. Route construction, 20 ants build tours in parallel using the GPU accelerated transition rule. Evaluation, the lengths of all 20 tours are calculated. The minimum length is tracked as the "Best Distance". Pheromone update, pheromones are evaporated and then reinforced on the edges of the best route. Model Training, the ML model is updated using the best route as the ground truth.

6. Evaluation

6.1 Spatial Topology Results

The spatial connectivity analysis provided quantitative confirmation of the clustered nature of last mile delivery. The analysis of 300,000 samples resulted in a connectivity graph with a density of approximately 0.00815.

The DBSCAN clustering identified 22 distinct operational zones¹⁶. The visualization of these clusters clearly shows how delivery points are grouped into neighbourhood's, separated by transit corridors.




Figure 1: Visualizing the spatial distribution of order locations. The distinct groupings correspond to residential and commercial districts in Shanghai.




Figure 2: The connectivity graph illustrates the network topology. Nodes with high connectivity degree represent central locations that serve as bridges between different delivery clusters.

6.2 Comparative Optimization Results

The three hybrid algorithms were evaluated on a test instance of 50 delivery locations over 100 iterations.

6.2.1 ANN-ACO Results

The ANN-ACO model showed inconsistent performance. In the final evaluation, it achieved a best distance of **3.79 km**. While it did optimize the route compared to the initial random state, the learning curve was erratic. The ANN struggled to generalize from the sparse feedback provided by the ants, often getting stuck in local optima where it would confidently predict poor edges as "good" simply because they were part of an early, suboptimal best route.

6.2.2 KNN-ACO Results

The KNN-ACO model demonstrated superior performance, achieving a final distance of **2.46 km**. This represents a 73.85% improvement over the initial state. The KNN-ACO model demonstrated the strongest performance among the proposed hybrid variants and achieved route lengths competitive with established heuristic baselines under the simplified experimental setting. The KNN model effectively acted as a "local expert". By recalling that moves in this direction worked well nearby, it guided the ants to complete clusters efficiently before moving to the next, minimizing cross town travel.




Figure 3: The convergence plot reveals the learning dynamics. The KNN-ACO (Gold) curve descends rapidly and stabilizes, indicating robust learning. The LSTM-ACO (Green) curve remains flat, indicating a failure to learn.

6.2.3 LSTM-ACO Results

The LSTM-ACO model yielded the poorest results, with a final distance of **7.88 km**. The sequence based prediction failed to provide meaningful guidance. In many cases, the LSTM guidance actively conflicted with the spatial heuristics, causing the ants to take erratic, zig zagging paths that drastically increased the total travel distance.

6.3 Comprehensive Performance Summary

The table below summarizes the final experimental results compared against the LaDe paper baselines.

Table 1: Final Performance Metrics

Method	Euclidean Distance (km)	Hit Rate @ Top 3 (%)
ANN-ACO	3.79	50.00%
KNN-ACO	2.56	68.46%
LSTM-ACO	7.88	51.92%

6.4 Statistical Comparison Across Multiple Routes

To ensure that the observed performance differences were not limited to a single illustrative route, each hybrid model was evaluated across 30 independently sampled routing instances. For each trial, a random subset of 50 delivery locations was selected and the corresponding optimization process was executed under identical parameter settings.

Table X reports the mean and standard deviation of the final route lengths obtained across all test routes. The KNN-ACO model consistently achieved the lowest average distance and exhibited the smallest variance, indicating both superior performance and higher stability. In contrast, the ANN-ACO model displayed moderate performance with higher variability, while the LSTM-ACO model produced significantly longer routes and showed the greatest dispersion. These results confirm that the advantage of the KNN based hybrid approach is robust across multiple problem instances and not attributable to a single favourable initialization.

Table 2: Mean and Standard Deviation of Route Length Across 20 Multiple Test Routes

Method	Mean Distance (km)	Std Distance (km)
ANN-ACO	2.15	0.25
KNN-ACO	3.12	0.83
LSTM-ACO	4.22	0.42




Figure 4: Bar chart comparing the final route distances and convergence. The superiority of the KNN-ACO approach is evident.




Figure 5: The resulting routes plotted on the map. The KNN-ACO route (Center) forms a coherent, circular path.

7. Discussion

7.1 The Supremacy of Local Spatial Learning

The most significant finding of this study is the outperformance of the KNN-based hybrid. This result aligns with the "First Law of Geography" everything is related to everything else, but near things are more related than distant things.

The VRP is inherently a problem of geometry. The quality of a move from A to B is heavily dependent on the local context. KNN explicitly captures this local context was efficient for a previous ant in the same neighbourhood, it is likely efficient for the current ant. The KNN model acts as a "spatial memory," allowing the swarm to share local navigational insights. This confirms the hypothesis that incorporating local spatial awareness is more valuable than global function approximation for this class of problems.

7.2 The Limitations of Deep Learning in VRP

Contrary to the trend of applying Deep Learning to every domain, this research highlights its limitations in specific combinatorial contexts. Data Scarcity which is Deep models like ANNs and LSTMs require massive amounts of data to converge. In this simulation, the training data is generated on the fly (one best route per epoch). This is insufficient for a deep network to learn a stable representation.

The Set vs Sequence Problem which is the failure of the LSTM is particularly illuminating. LSTMs excel at sequence problems (like language) where the order is intrinsic to the meaning. However, a set of delivery points is a set, not a sequence. The optimal order is an emergent property of the geometry, not a latent property of the input history. By forcing a sequential interpretation on a geometric problem, the LSTM introduced noise rather than signal.

7.3 Practical Implications

For the logistics industry, these findings are highly relevant. They suggest that computationally expensive "black box" deep learning models are not always necessary for route optimization. A hybrid system combining robust metaheuristics (ACO) with simple, interpretable memory based learning (KNN) can provide high quality solutions with a fraction of the training overhead. This is particularly valuable for edge computing scenarios where delivery trucks might need to re-optimize routes using onboard hardware with limited processing power.

8. Conclusion and Future Work

8.1 Conclusion

This research project successfully developed and evaluated a Hybrid Ant Colony Optimization framework for large scale last mile delivery, utilizing the real world LaDe dataset. The study rigorously tested the hypothesis that machine learning guidance can enhance swarm intelligence.

The results conclusively demonstrate that KNN-ACO is the most effective architecture among those tested. It achieved a 73.85% improvement over random routing and produced solutions competitive with industry baselines. The success of the KNN approach underscores the importance of spatial locality in routing problems. Conversely, the poor performance of the LSTM approach serves as a cautionary tale against the indiscriminate application of deep

learning architectures to combinatorial problems without regard for the problem's underlying mathematical structure.

8.2 Future Work

1. **Graph Neural Networks (GNN):** Since the delivery network is structured as a graph, GNNs represent a more theoretically sound deep learning approach than ANNs or LSTMs. Future work should explore GNN-ACO hybrids, where the GNN learns node embeddings that capture the topology of the road network.
2. **Multi Objective Optimization:** Real world delivery involves trading off distance against time and vehicle capacity. Expanding the cost function to a Multi Objective formulation would increase the practical applicability of the system.
3. **Dynamic Replanning:** The LaDe dataset includes timestamps. Future research could simulate a dynamic environment where new orders arrive in real time, testing the hybrid system's ability to adapt its routes on the fly.

References

1. Ahmed, Z.H., Hameed, A.S. and Mutar, M.L. (2023) 'An Enhanced Ant Colony System Algorithm Based on Subpaths for Solving the Capacitated Vehicle Routing Problem', *ResearchGate*.
2. Bai, R., et al. (2023) 'Integrating Machine Learning Into Vehicle Routing Problem: Methods and Applications', *IEEE Access*.
3. Bell, J.E. and McMullen, P.R. (2004) 'Ant colony optimization techniques for the vehicle routing problem', *Advanced Engineering Informatics*, 18(1), pp. 41-48.
4. Bengio, Y., Lodi, A. and Prouvost, A. (2021) 'Machine learning for combinatorial optimization: a methodological tour d'horizon', *European Journal of Operational Research*, 290(2), pp. 405-421.
5. Dantzig, G.B. and Ramser, J.H. (1959) 'The truck dispatching problem', *Management Science*, 6(1), pp. 80-91.
6. Dorigo, M. (1992) *Optimization, Learning and Natural Algorithms*. PhD Thesis, Politecnico di Milano.
7. Dorigo, M. and Stützle, T. (2004) *Ant Colony Optimization*. Cambridge, MA: MIT Press.
8. Foa, S., Coppola, C., Grani, G. and Palagi, L. (2023) 'Deep Reinforcement Learning for Solving Multi-objective Vehicle Routing Problem', *arXiv preprint*.
9. Gambardella, L.M. and Dorigo, M. (1999) 'HAS-SOP: An hybrid ant system for the sequential ordering problem', *Journal of Global Optimization*, 14, pp. 265–285.
10. Karagul, K. et al. (2024) 'Machine Learning for Enhancing Metaheuristics in Global Optimization: A Comprehensive Review', *Mathematics*, 13(18).
11. Kool, W., van Hoof, H. and Welling, M. (2019) 'Attention, Learn to Solve Routing Problems!', *International Conference on Learning Representations (ICLR)*.
12. Laporte, G. (2009) 'Fifty years of vehicle routing', *Transportation Science*, 43(4), pp. 408-416.
13. Li, F. et al. (2021) 'Deep Reinforcement Learning for Routing Problems: A Comprehensive Review', *IEEE Access*, 9.
14. Mahi, H., Baykan, O.K. and Kodaz, N. (2015) 'A new hybrid method based on particle swarm optimization, ant colony algorithm and 3-opt algorithm for traveling salesman problem', *Applied Soft Computing*, 30, pp. 484-490.
15. Nazari, M. et al. (2018) 'Reinforcement Learning for Solving the Vehicle Routing Problem', *Advances in Neural Information Processing Systems*.
16. Toth, P. and Vigo, D. (2014) *Vehicle Routing: Problems, Methods, and Applications*. Philadelphia: SIAM.

17. Vinyals, O., Fortunato, M. and Jaitly, N. (2015) 'Pointer Networks', *Advances in Neural Information Processing Systems (NIPS)*.
18. Wang, L. et al. (2023) 'LaDe: The First Comprehensive Last-mile Delivery Dataset from Industry', *arXiv preprint arXiv:2306.10675*.
19. Wu, L., Wen, H. and Hu, H. (2023) 'LaDe: The First Comprehensive Last-mile Delivery Dataset from Industry', *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*.
20. Zhang, J. et al. (2020) 'Deep Reinforcement Learning for Combinatorial Optimization: A Survey', *IEEE Transactions on Cybernetics*.
21. Zhang, X. et al. (2023) 'A multi-agent deep reinforcement learning approach for solving the multi-depot vehicle routing problem', *Journal of Management Analytics*, 10(9).